

Project Overview

This repository contains data analysis and visualization using Jupiter Notebook And Tableau to explore the relationship between housing market indicators and key factors that influence home prices nationally.

Data Collection

The dataset utilized for this analysis comprises the following housing market indicators

Indicator : Source

Label (HPI) : <https://fred.stlouisfed.org/series/CSUSHPIA>

GDP : <https://fred.stlouisfed.org/graph/?g=znfe>

UNEMPLOYMENT RATE : <https://fred.stlouisfed.org/series/UNRATE>

MORTGAGE RATE : <https://fred.stlouisfed.org/graph/?g=zneW>

Inflation : <https://www.usinflationcalculator.com/inflation/current-inflation-rates/>

Population growth rate : <https://fred.stlouisfed.org/series/POPTHM>

Monthly supply rate : <https://fred.stlouisfed.org/series/MSACSR>

Median Sales Price : <https://fred.stlouisfed.org/series/MSPNHSUS>

Housing Units Started: <https://fred.stlouisfed.org/series/HOUST1F>

Disposable Personal Income : <https://fred.stlouisfed.org/series/DSPI>

Personal Saving Rate : <https://fred.stlouisfed.org/series/PSAVERT>

Personal Consumption Expenditures : <https://fred.stlouisfed.org/series/PCE>

Total Construction Spending : <https://fred.stlouisfed.org/series/TTLCONS>

Data Analysis and Visualization

Problem Statement

Q. Find publicly available data for key factors that influence US home prices nationally. Then, build a data science model that explains how these factors impacted home prices over the last 20 years. Use the S&P Case-Schiller Home Price Index as a proxy for home prices: fred.stlouisfed.org/series/CSUSHPISA.

Importing libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.ensemble import RandomForestRegressor

import warnings
warnings.filterwarnings('ignore')

pd.options.display.max_columns = None
```

```
In [2]: df = pd.read_excel(r"C:\Users\kiran\Desktop\data\trained\Home LLC Interview\data\df")
```

Out[2]:

	Date	HPI	UNEMP_RATE	DSPI	GDP	HOUST1F	Inflation	M_SP	Supply_Rate
0	2003-01-01	128.461	5.8	8268.0	98.336393	1537	2.6	181700	4.0
1	2003-02-01	129.355	5.9	8274.7	98.329348	1301	3.0	187000	4.0
2	2003-03-01	130.148	5.9	8313.4	98.368090	1399	3.0	185100	4.0
3	2003-04-01	130.884	6.0	8342.9	98.453754	1374	2.2	189500	4.0
4	2003-05-01	131.735	6.1	8394.8	98.584413	1391	2.1	195500	4.0
...
244	2023-05-01	302.566	3.7	20185.4	99.943470	1012	4.0	421200	7.0
245	2023-06-01	304.593	3.6	20208.4	100.045476	930	3.0	417600	7.0
246	2023-07-01	306.767	3.5	20229.6	100.164368	988	3.2	435800	7.0
247	2023-08-01	309.155	3.8	20306.7	100.289174	948	3.7	439900	7.0
248	2023-09-01	311.175	3.8	20388.3	100.410037	968	3.7	422300	7.0

249 rows × 15 columns

```
In [3]: df.columns
```

Out[3]: Index(['Date', 'HPI', 'UNEMP_RATE', 'DSPI', 'GDP', 'HOUST1F', 'Inflation', 'M_SP', 'Supply_Rate', 'MORTGAGE', 'PCE', 'PSAVERT', 'POPTHM', 'Growth Rate', 'TTLCONS'], dtype='object')

Columns

- Date - The date of observation
- UNEMP_RATE - Unemployment Rate (in %)
- DSPI - Disposable Personal Income (in Billions)
- GDP - Gross Domestic Product (GDP): Normalised for United States
- HOUST1F - New Privately-Owned Housing Units Started (in Thousands)
- Inflation - U.S Annual Inflation Rates.
- M_SP - Median Sales Price for New Houses Sold in the United States
- Supply_Rate - Monthly Supply of New Houses in the United States
- MORTGAGE - Mortgage Interest Rate (in %)

- PCE - Personal Consumption Expenditures (in Billions)
- PSAVERT - Personal Saving Rate (in %)
- POPTHM - Population (in Thousands)
- Growth Rate - Population Growth Rate
- TTLCONS - Total Construction Spending (in Millions)

Target

1. HPI - S&P/Case-Shiller U.S. National Home Price Index

Analising basic metrics

In [4]: `df.shape`

Out[4]: (249, 15)

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 249 entries, 0 to 248
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date            249 non-null   datetime64[ns]
1   HPI             249 non-null   float64
2   UNEMP_RATE     249 non-null   float64
3   DSPi           249 non-null   float64
4   GDP            249 non-null   float64
5   HOUST1F        249 non-null   int64
6   Inflation      249 non-null   float64
7   M_SP          249 non-null   int64
8   Supply_Rate    249 non-null   float64
9   MORTGAGE       249 non-null   float64
10  PCE            249 non-null   float64
11  PSAVERT        249 non-null   float64
12  POPTHM         249 non-null   int64
13  Growth Rate    249 non-null   float64
14  TTLCONS        249 non-null   int64
dtypes: datetime64[ns](1), float64(10), int64(4)
memory usage: 29.3 KB
```

Findings

- There are 13 Numerical columns and 1 Date columns
- Date column must be converted into datetime data type

```
In [6]: # Convert the 'date_column' to datetime data type
df['Date'] = pd.to_datetime(df['Date'])

# Extract Month and Year
df['month'] = df['Date'].dt.month
df['year'] = df['Date'].dt.year
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: Date          0
        HPI           0
        UNEMP_RATE    0
        DSPI          0
        GDP           0
        HOUST1F       0
        Inflation     0
        M_SP          0
        Supply_Rate   0
        MORTGAGE      0
        PCE           0
        PSAVERT       0
        POPTHM        0
        Growth Rate   0
        TTLCONS       0
        month         0
        year          0
        dtype: int64
```

- No Null Values can be observed

```
In [8]: df.nunique()
```

```
Out[8]: Date          249
        HPI           249
        UNEMP_RATE     64
        DSPI          248
        GDP           249
        HOUST1F       220
        Inflation      72
        M_SP          227
        Supply_Rate    72
        MORTGAGE      238
        PCE           249
        PSAVERT        78
        POPTHM        249
        Growth Rate    249
        TTLCONS       249
        month         12
        year          21
        dtype: int64
```

- All columns are continuous

```
In [9]: df.describe()
```

Out[9]:

	HPI	UNEMP_RATE	DSPI	GDP	HOUST1F	Inflation	M_SP
count	249.000000	249.000000	249.000000	249.000000	249.000000	249.000000	249.00
mean	185.081968	5.924498	13179.864659	99.870250	917.212851	2.553414	283452.61
std	46.531700	2.047836	3280.418809	1.215264	389.720928	1.910771	70742.59
min	128.461000	3.400000	8268.000000	92.031035	353.000000	-2.100000	181700.00
25%	148.278000	4.400000	10745.400000	99.622947	614.000000	1.500000	228300.00
50%	174.342000	5.400000	12441.900000	99.943470	843.000000	2.200000	262200.00
75%	202.913000	7.200000	15508.800000	100.421834	1136.000000	3.400000	321400.00
max	311.175000	14.700000	21858.100000	101.871442	1823.000000	9.100000	496800.00

Observations

- Distribution
 - Majority of the data seems Normally Distributed (Since Mean \approx Median), But a hint of skewness can be observed in few columns.
 - Normally distributed (Since Mean \approx Median)
 - HPI
 - UNEMP_RATE
 - DSPI
 - GDP
 - HOUST1F
 - Inflation
 - M_SP
 - Supply_Rate
 - MORTGAGE
 - PCE
 - PSAVERT
 - POPTHM
 - Growth Rate
 - TTLCONS
 - Right skewed (mean > median)
 - UNEMP_RATE
 - PSAVERT
 - LEFT skewed (mean < median)
 - GDP
 - Growth Rate
- From Above observation we cannot find the presence of outliers(i.e. No Differences Between Max & 75%,Min & 25%) can be observed

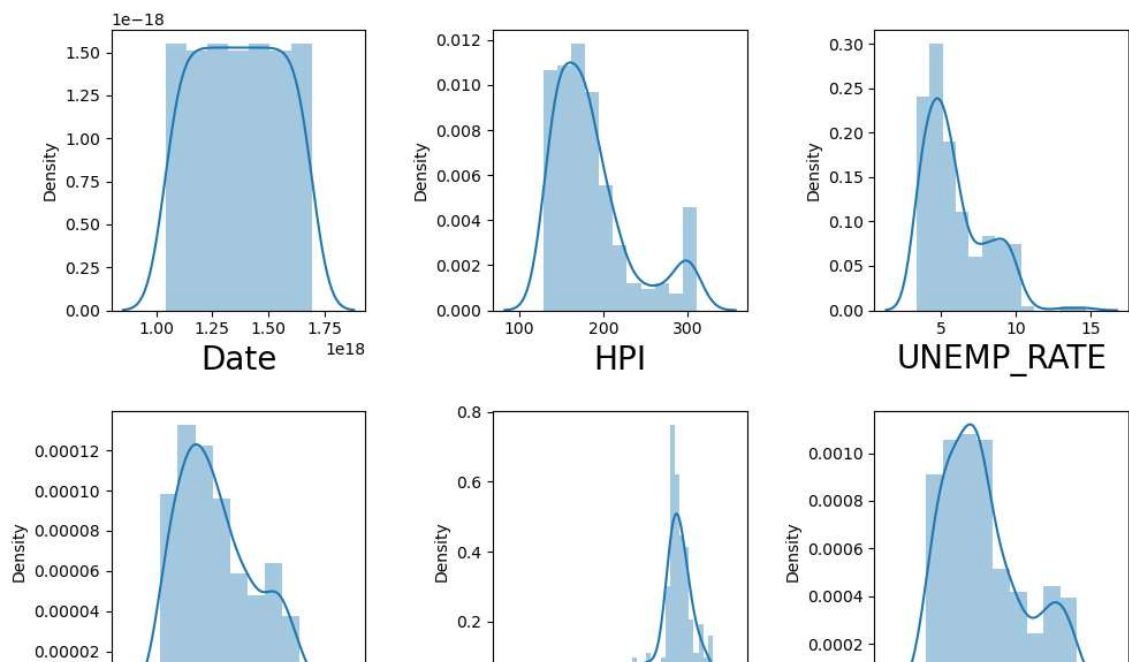
Data Visualization

Distribution Plot

```
In [10]: plt.figure(figsize=(10,20))
pn = 1
for column in df.columns:

    sns.distplot(df[column],ax=plt.subplot(6,3,pn))
    plt.xlabel(column,fontsize = 20)
    pn +=1

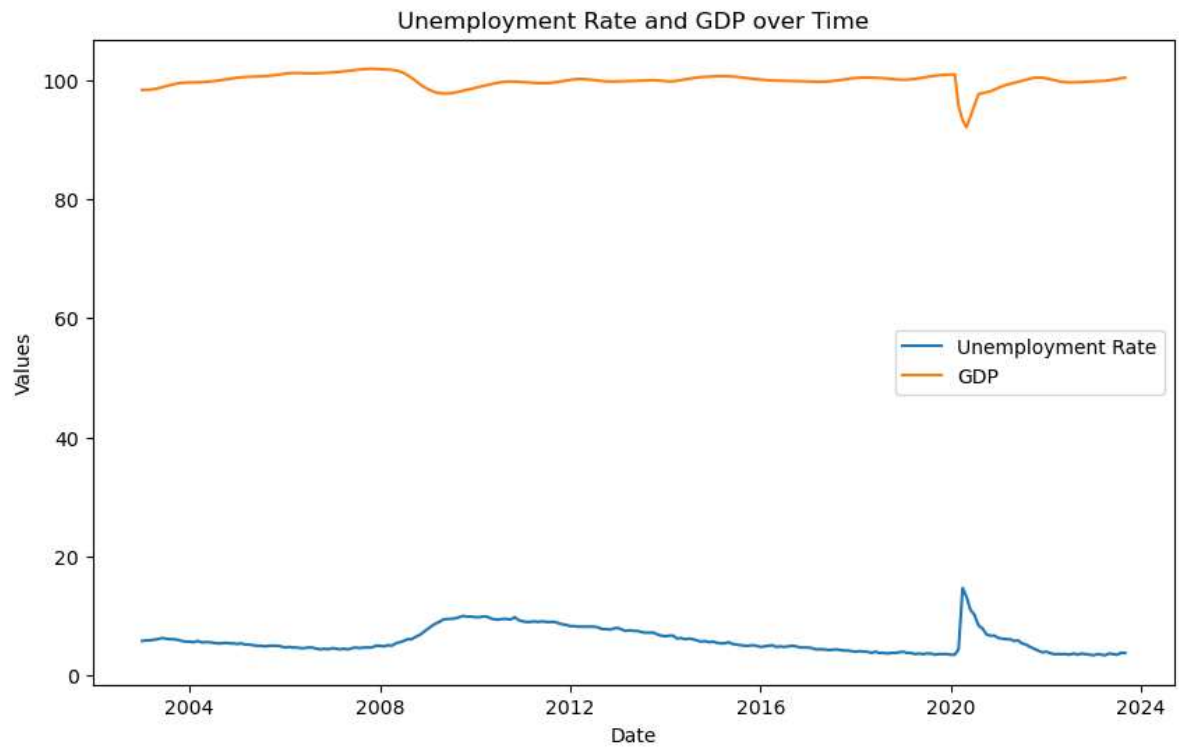
plt.tight_layout()
```



- We can observe that Most of the columns are Not skewed (Normally Distrubuted)
- Skeweness is as discussed earlier

Line Plot

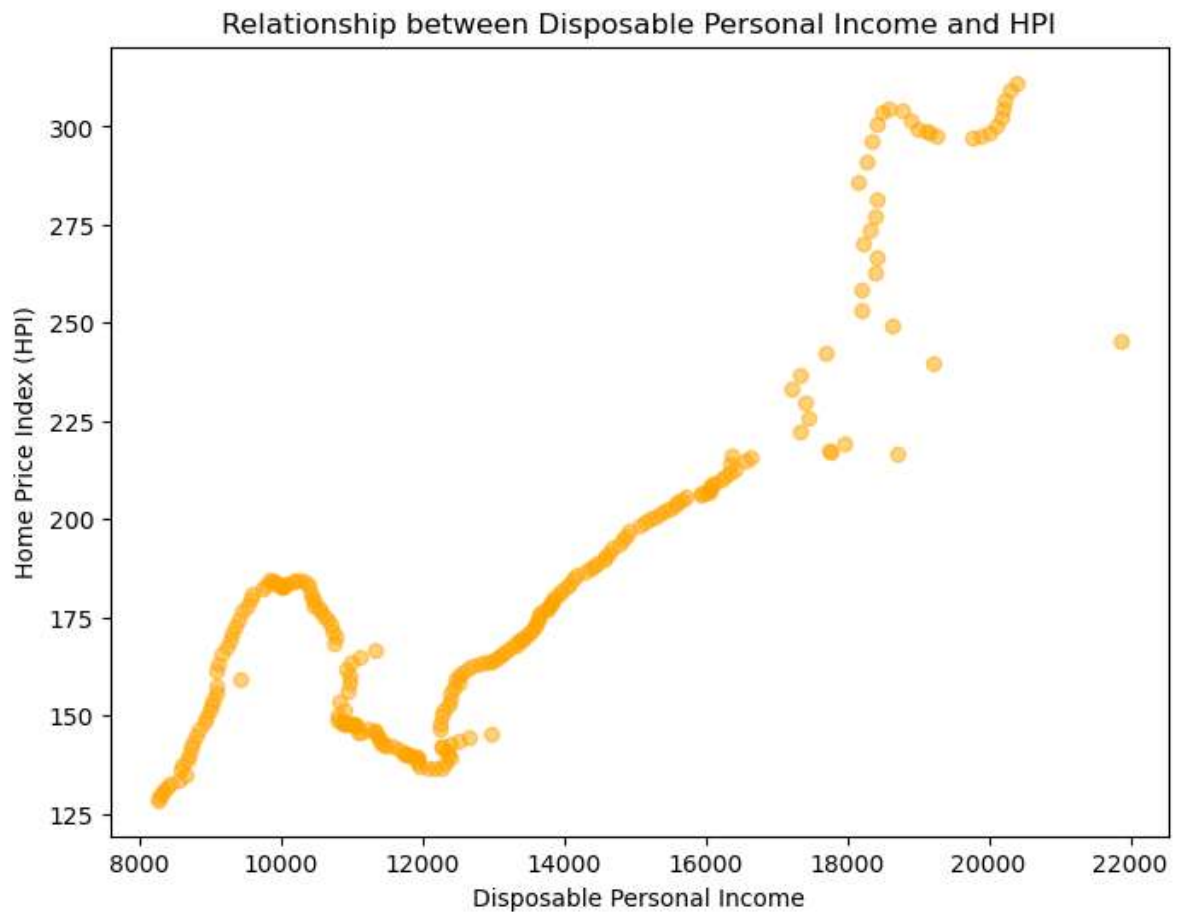
```
In [11]: plt.figure(figsize=(10, 6))
plt.plot(df['Date'], df['UNEMP_RATE'], label='Unemployment Rate')
plt.plot(df['Date'], df['GDP'], label='GDP')
plt.xlabel('Date')
plt.ylabel('Values')
plt.title('Unemployment Rate and GDP over Time')
plt.legend()
plt.show()
```



- we can observe that GDP and Unemployment rate are inversly related i.e. increase in unemploymnt rate decreases GDP

Scatter Plot


```
In [12]: plt.figure(figsize=(8, 6))
plt.scatter(df['DSPI'], df['HPI'], alpha=0.5, color='orange')
plt.xlabel('Disposable Personal Income')
plt.ylabel('Home Price Index (HPI)')
plt.title('Relationship between Disposable Personal Income and HPI')
plt.show()
```




- we can observe that HPI AND Disposable Personal Income are **positively** correlated

Year wise Median sale price

```
In [13]: median_price_by_year = df.groupby('year')['M_SP'].median()
plt.figure(figsize=(10, 6))
bar_plot = median_price_by_year.plot(kind='bar', color='skyblue')


for idx, value in enumerate(median_price_by_year):
    plt.text(idx, value + 1000, f'{value:.0f}', ha='center', va='bottom', font

plt.xlabel('Year')
plt.ylabel('Median Sales Price')
plt.title('Median Sales Price by Year')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Corelation Analysis

```
In [14]: plt.figure(figsize=(20,20))
sns.heatmap(df.corr(),annot=True,annot_kws={'size':10})
plt.tight_layout()
```



Important Observations

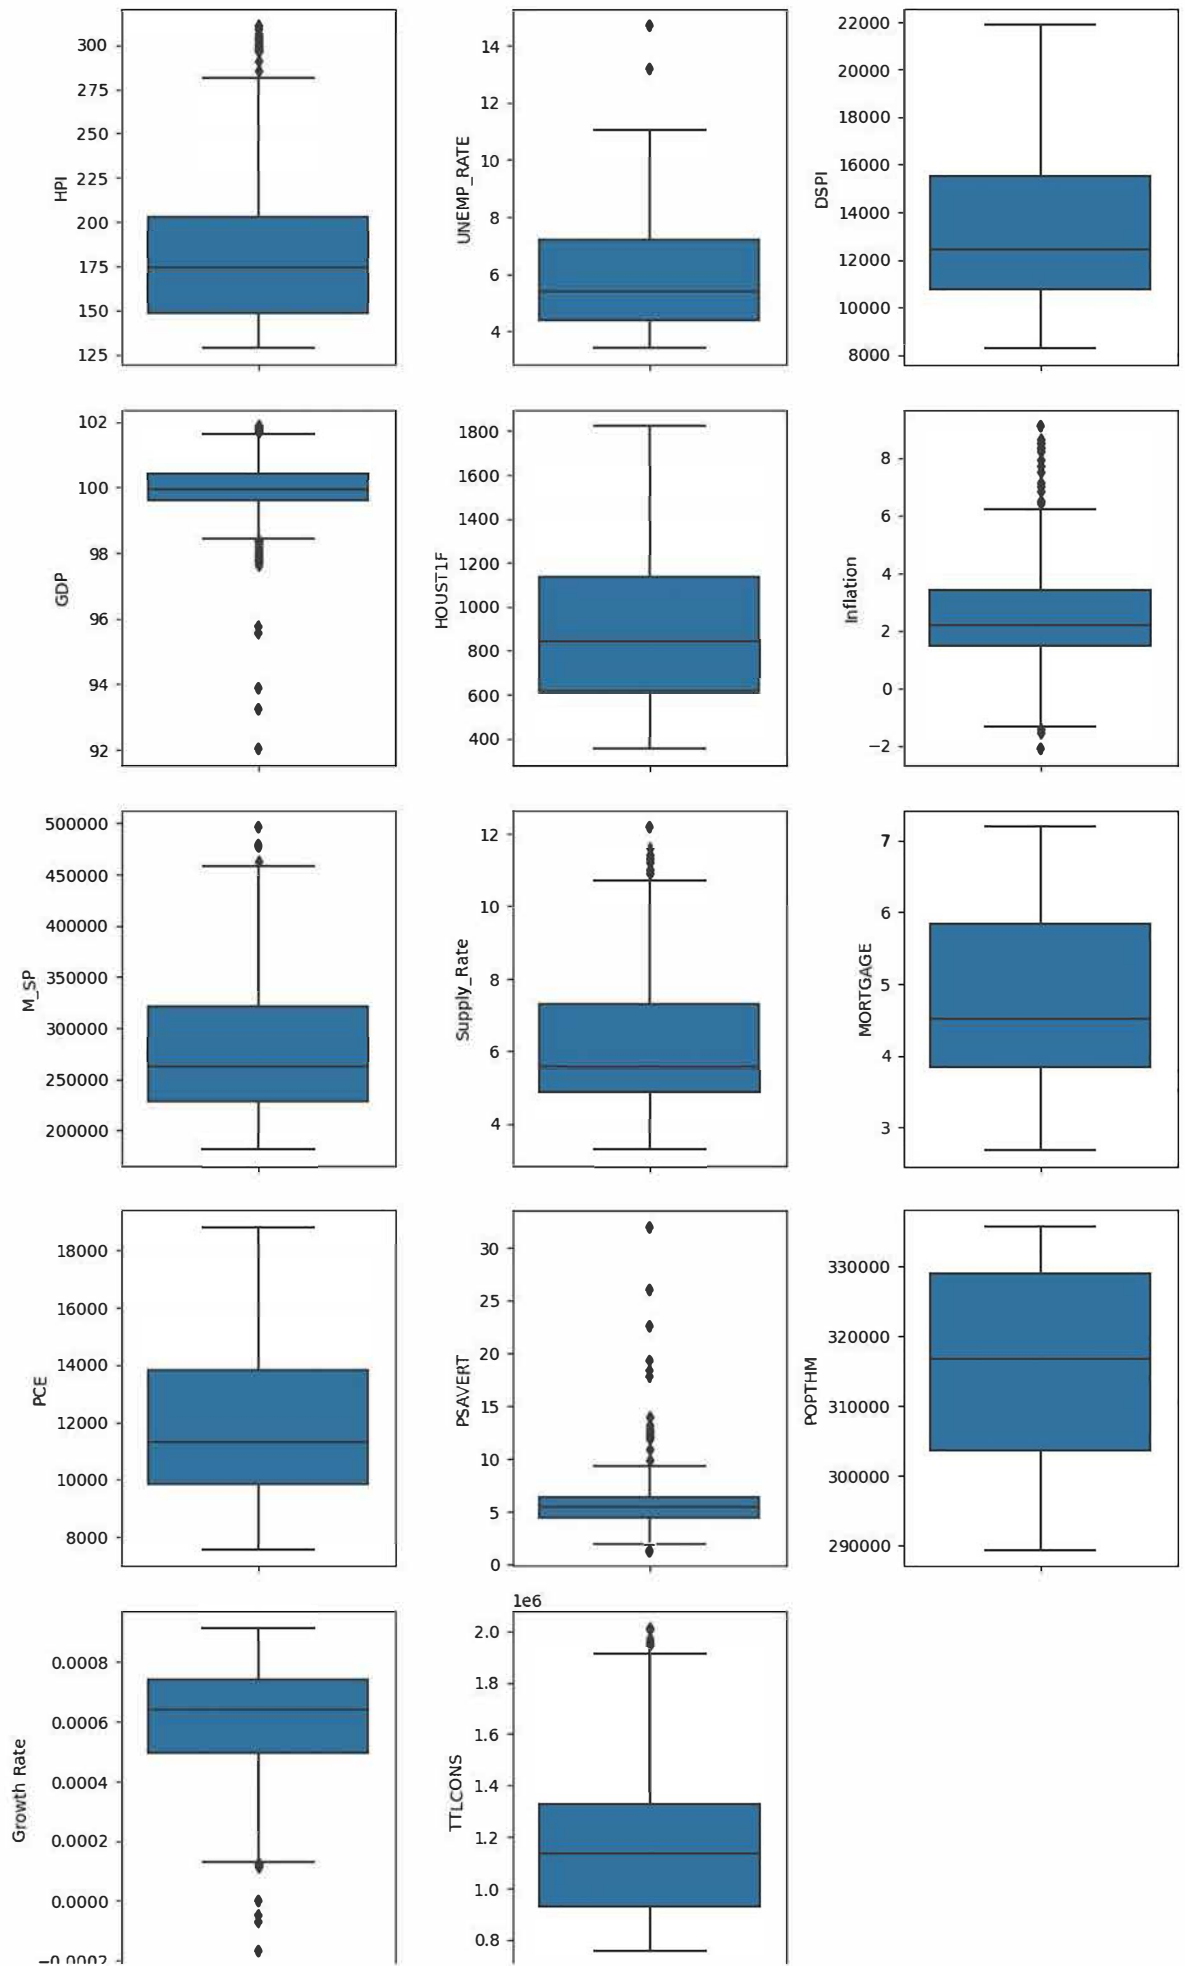
- HPI:
 - Strong positive correlations with: DSPI (0.858), M_SP (0.941), PCE (0.886), TTLCONS (0.983)
- UNEMP_RATE:
 - negative correlations with: GDP (-0.617), HOUST1F (-0.471), PSAVERT (-0.429)
- DSPI:
 - Strong positive correlations with: HPI (0.858), M_SP (0.930), PCE (0.982), TTLCONS (0.843)
- GDP:

- Strong negative correlation with: UNEMP_RATE (-0.617)
- Inflation:
 - Moderate positive correlations with: HPI (0.597), M_SP (0.494), PCE (0.388), TTLCONS (0.538)
- M_SP:
 - Strong positive correlations with: HPI (0.941), DSPI (0.930), PCE (0.958), TTLCONS (0.921)
- Growth Rate:
 - Strong negative correlations with: DSPI (-0.824), TTLCONS (-0.656), year (-0.806)
- POPTHM:
 - Strong positive correlations with: DSPI (0.948), HPI (0.701), year (0.992)

Outliers

```
In [15]: df1 = df.drop(columns=['month', 'year', 'Date']).copy()
```

```
In [16]: plt.figure(figsize=(10,20))
pn=1
for i in df1.columns:
    sns.boxplot(data=df1,y=i,ax=plt.subplot(6,3,pn))
    pn+=1
plt.tight_layout()
```

- There are Many outliers which are needed to be Checked


```
In [17]: # finding outliers using z score
from scipy.stats import zscore
z = np.abs(zscore(df1))
z[z>3].count()
```

```
Out[17]: HPI          0
          UNEMP_RATE  2
          DSPI        0
          GDP         5
          HOUST1F     0
          Inflation   6
          M_SP        1
          Supply_Rate 1
          MORTGAGE    0
          PCE         0
          PSAVERT     6
          POPTHM      0
          Growth Rate  3
          TTLCONS     0
          dtype: int64
```



```
In [18]: sns.heatmap(z[z>3])
```

```
Out[18]: <AxesSubplot: >
```



- We can observe that Majority of the outliers are around years 2021,2022.
- This can be due to **Covid 19 Pandemic**
- so i,Feel It is wise to keep these outliers for gaining valuable insights.

Treating Skewness

```
In [19]: df1.skew().abs()
```

```
Out[19]: HPI          1.311186
          UNEMP_RATE  1.051399
          DSPI        0.556285
          GDP         2.543659
          HOUST1F     0.726292
          Inflation   1.082925
          M_SP        0.993782
          Supply_Rate 0.932554
          MORTGAGE    0.197052
          PCE         0.645530
          PSAVERT     3.648431
          POPTHM      0.266724
          Growth Rate  1.047579
          TTLCONS     0.912868
          dtype: float64
```

```
In [20]: skewed_columns = [x for x in df1.columns if df1[x].skew() > 1 or df1[x].skew()
          skewed_columns
```

```
Out[20]: ['HPI', 'UNEMP_RATE', 'GDP', 'Inflation', 'PSAVERT', 'Growth Rate']
```

skewness of 'HPI', 'UNEMP_RATE', 'GDP', 'Inflation', 'PSAVERT' & 'Growth Rate' Needs to be Treated

```
In [21]: from sklearn.preprocessing import PowerTransformer
          pt = PowerTransformer(method='yeo-johnson', standardize=False)
          df1[skewed_columns] = pt.fit_transform(df1[skewed_columns])
          df1
```

```
Out[21]:
```

	HPI	UNEMP_RATE	DSPI	GDP	HOUST1F	Inflation	M_SP	Supply_Rate
0	0.539097	0.875639	8268.0	3.787013e+61	1537	2.016832	181700	4.0
1	0.539098	0.877936	8274.7	3.778539e+61	1301	2.271700	187000	4.5
2	0.539099	0.877936	8313.4	3.825366e+61	1399	2.271700	185100	4.1
3	0.539099	0.880169	8342.9	3.930911e+61	1374	1.752193	189500	4.1
4	0.539100	0.882340	8394.8	4.097343e+61	1391	1.684320	195500	3.9
...
244	0.539149	0.805378	20185.4	6.286895e+61	1012	2.873938	421200	7.2
245	0.539149	0.800467	20208.4	6.490686e+61	930	2.271700	417600	7.5
246	0.539149	0.795343	20229.6	6.736289e+61	988	2.395884	435800	7.1
247	0.539149	0.810088	20306.7	7.003786e+61	948	2.697932	439900	7.8
248	0.539150	0.810088	20388.3	7.272619e+61	968	2.697932	422300	7.2

249 rows × 14 columns



Removing Unwanted Columns

- The Following columns have High Multicollinearity and Presence of direct High correlation with HPI
 - M_SP
 - TTLCONS
- For Better Model Training And to understand Feature Importance Better, We can remove these columns

In [22]:

```
x = df1.drop(columns=['HPI', "M_SP", "PCE", "TTLCONS "])
y=df1['HPI']
```

Standard Scaling

In [23]: **from** sklearn.preprocessing **import** StandardScaler

```
scaler = StandardScaler()
x_sc = scaler.fit_transform(x)
```

Model Building

Linear Regression

In [24]:

```
# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(x_sc, y, test_size=0.2, ra

# Initializing and fitting the linear regression model
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)

# Predicting on the test set
y_pred = linear_reg.predict(X_test)

# Evaluating the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Linear Regression Mean Squared Error: {mse}")
print(f"Linear Regression R-squared: {r2}")
```

Linear Regression Mean Squared Error: 7.421402640693978e-12

Linear Regression R-squared: 0.9529587278313972

Random Forest Regression

In [25]:

```
# Initializing and fitting the random forest regression model
rf_reg = RandomForestRegressor(random_state=42)
rf_reg.fit(X_train, y_train)

# Predicting on the test set
y_pred_rf = rf_reg.predict(X_test)

# Evaluating the model
mse_rf = mean_squared_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)

print(f"Random Forest Regression Mean Squared Error: {mse_rf}")
print(f"Random Forest Regression R-squared: {r2_rf}")
```

Random Forest Regression Mean Squared Error: 6.872577556948824e-13

Random Forest Regression R-squared: 0.9956437508243586

- We can observe that our Models Linear regression and Random Forest classifier Have Given Good R2 Scores i.e. 95% and 99% respectively

Getting Features or coefficient details from models

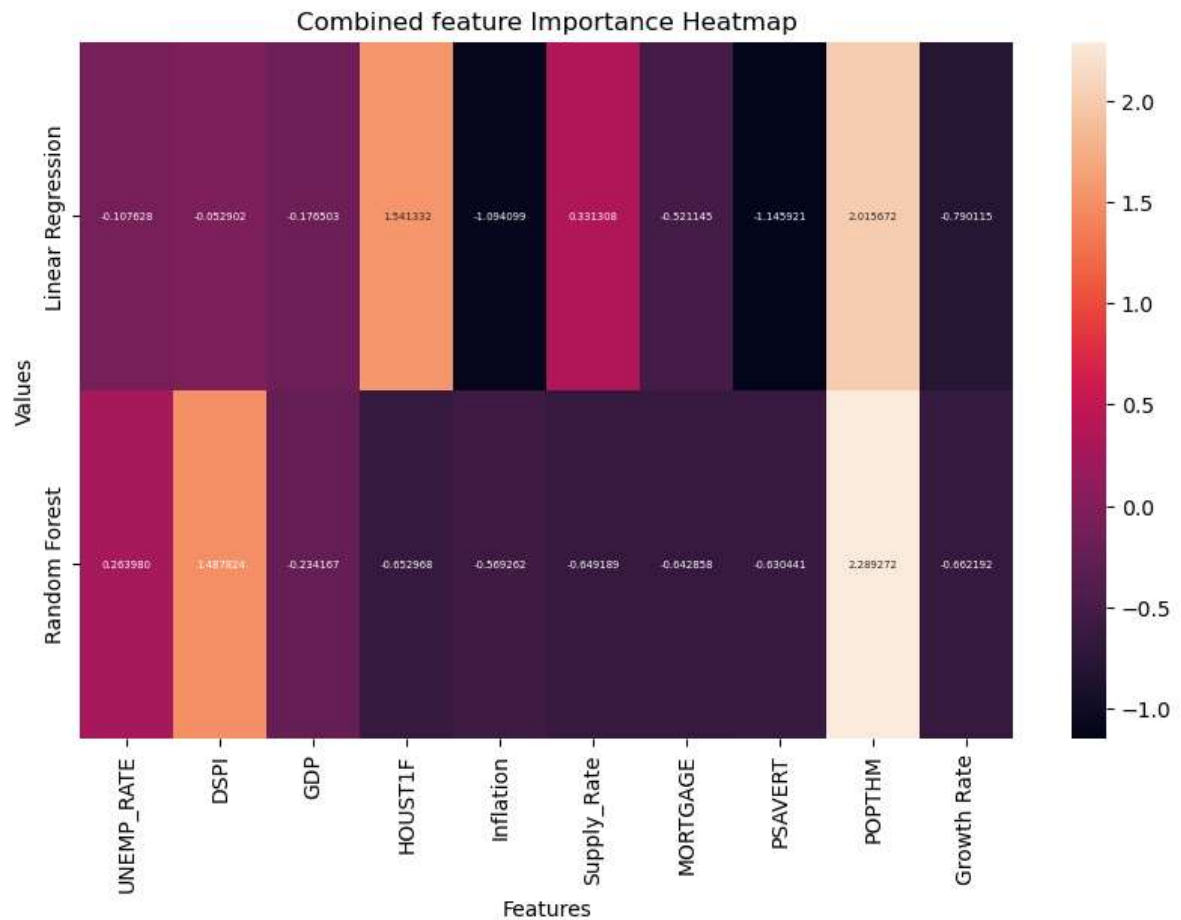
In [26]:

```
# Coefficients for Linear Regression
linear_coefficients = linear_reg.coef_
sc_l_coef = scaler.fit_transform(linear_coefficients.reshape(10,1))

# Feature importance for Random Forest
rf_feature_importances = rf_reg.feature_importances_
sc_rf_coef = scaler.fit_transform(rf_feature_importances.reshape(10,1))
```

Plotting Heatmaps of feature Importance

```
In [27]: plt.figure(figsize=(10, 6))
sns.heatmap(np.reshape(list(sc_l_coef) + list(sc_rf_coef),(2,10)), annot=True,
plt.title('Combined feature Importance Heatmap')
plt.xlabel('Features')
plt.ylabel('Values')
plt.show()
```



observations

- FROM both Models we can observe that the important factors affecting Housing price index are
 - Population
 - Disposable Personal Income
 - inflation
 - Personal Saving Rate
 - Housing Units Started

Test Prediction

lets predict each model on trail data by changing 1 variable by 1 unit in each step

```
In [28]: z = np.array([[0,0,0,0,0,0,0,0,0,0],
                        [1,0,0,0,0,0,0,0,0,0],
                        [0,1,0,0,0,0,0,0,0,0],
                        [0,0,1,0,0,0,0,0,0,0],
                        [0,0,0,1,0,0,0,0,0,0],
                        [0,0,0,0,1,0,0,0,0,0],
                        [0,0,0,0,0,1,0,0,0,0],
                        [0,0,0,0,0,0,1,0,0,0],
                        [0,0,0,0,0,0,0,1,0,0],
                        [0,0,0,0,0,0,0,0,1,0],
                        [0,0,0,0,0,0,0,0,0,1]])
z.shape
```

```
Out[28]: (11, 10)
```

```
In [29]: z_pred_rf = rf_reg.predict(z)
z_pred_rf
```

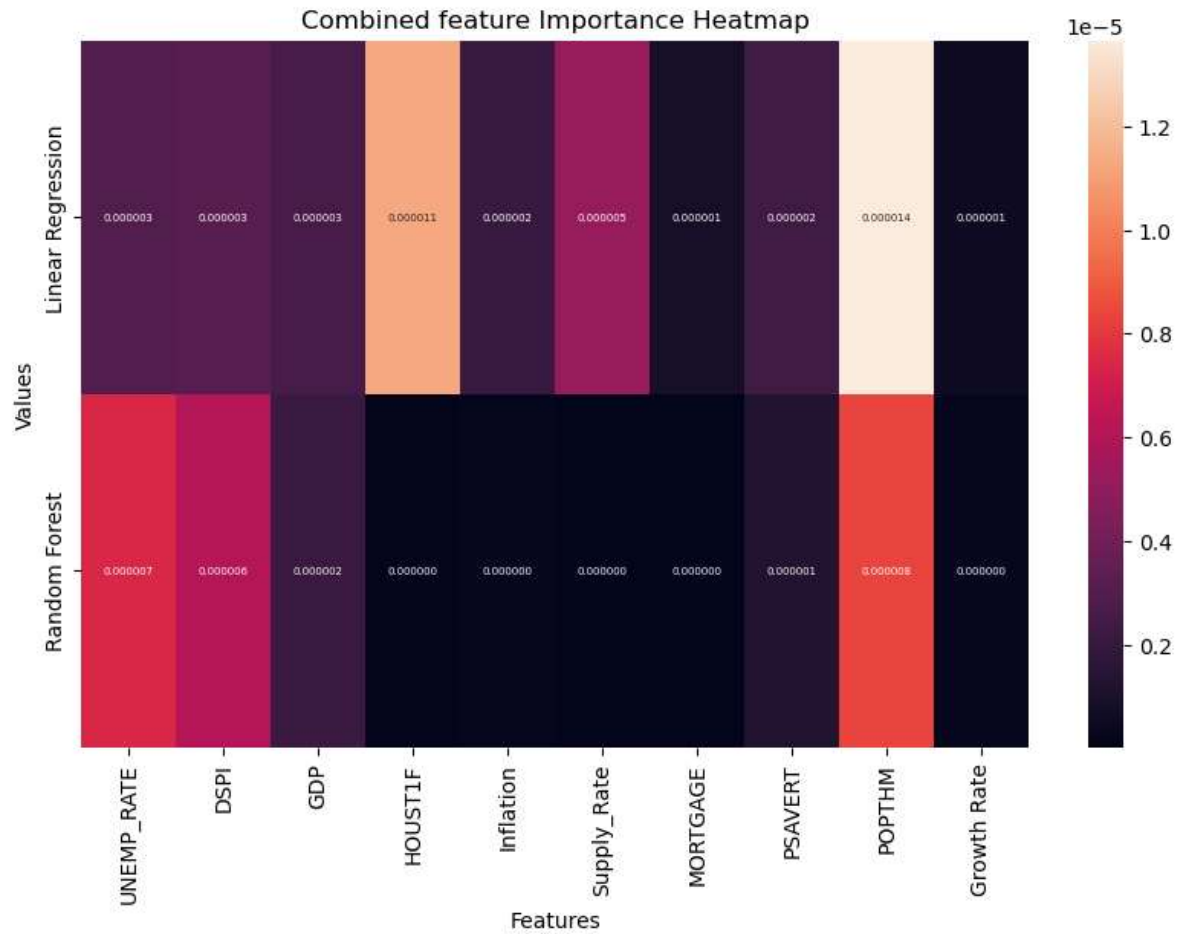
```
Out[29]: array([0.53911969, 0.53911224, 0.53912575, 0.53912182, 0.53911958,
                0.53911949, 0.53911968, 0.53911964, 0.53911843, 0.53912804,
                0.53911943])
```

```
In [30]: z_pred_lr = linear_reg.predict(z)
z_pred_lr
```

```
Out[30]: array([0.53912467, 0.53912759, 0.53912787, 0.53912725, 0.53913593,
                0.53912261, 0.53912981, 0.5391255 , 0.53912235, 0.53913833,
                0.53912414])
```

```
In [31]: Feature_Importance = np.reshape(list(np.abs(z_pred_lr[1:] - z_pred_lr[0])) + li
```

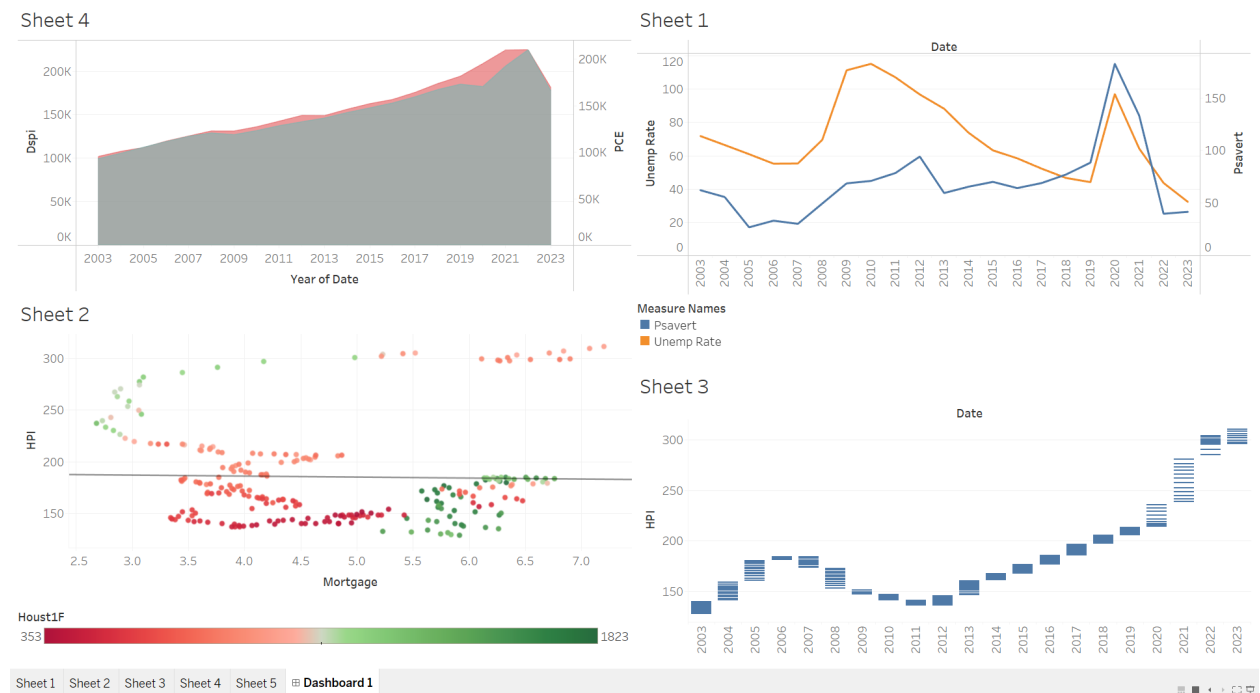
```
In [32]: plt.figure(figsize=(10, 6))
sns.heatmap(Feature_Importance, annot=True,fmt='.6f',annot_kws={'size':5}, yticklabels=
plt.title('Combined feature Importance Heatmap')
plt.xlabel('Features')
plt.ylabel('Values')
plt.show()
```



Conclusion

- Primary Factor affecting Housing Price index is **Population**
- Other Factors that impact HPI are *Unemployment Rate and Housing Units Started*

Power BI Dashboard



1. Sheet 1
 - a. We can observe that unemployment Rate and Public saving rate are positively co related
 - b. Effect of Covid 19 can be clearly Observed in years of 2019,2020,2021,2022
2. Sheet 2
 - a. New House start Rate is High when mortgage rates are High and HPI is low
 - b. New House start Rate is low when mortgage rates are Low and HPI is low
3. Sheet 3
 - a. Due to Covid19 Pandemic we can see variation in inflation over years of 2021 and 2022
4. Sheet 4
 - a. High Positive correlation can be observed between Disposable Personal Income and Personal Consumption Expenditures